

preprocessing

November 3, 2024

```
[123]: # -*- coding: utf-8 -*-  
#  
# Licensed under the Apache License, Version 2.0 (the "License");  
# you may not use this file except in compliance with the License.  
# You may obtain a copy of the License at  
#  
# http://www.apache.org/licenses/LICENSE-2.0  
#  
# Unless required by applicable law or agreed to in writing, software  
# distributed under the License is distributed on an "AS IS" BASIS,  
# WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or  
# implied.  
# See the License for the specific language governing permissions and  
# limitations under the License.  
#
```

0.0.1 Adam Candrák/Mária Matušisková - 50%/50%

1 Imports

```
[124]: import pandas as pd  
import numpy as np  
import seaborn as sns  
from sklearn.feature_selection import SelectKBest, f_classif  
import matplotlib.pyplot as plt  
from numpy import mean  
from numpy import std  
from collections import Counter  
from sklearn.model_selection import train_test_split  
from sklearn.pipeline import Pipeline  
from sklearn.compose import ColumnTransformer  
from sklearn.preprocessing import StandardScaler, PowerTransformer, MinMaxScaler
```

1.1 Global variables

```
[125]: target_column = 'mwra'
test_size = 0.3
random_state = 42
```

```
[126]: connections_file = "../data/Connections.csv"
processes_file = "../data/Processes.csv"

connections = pd.read_csv(connections_file, sep='\t')
processes = pd.read_csv(processes_file, sep='\t')
```

Change the names of the columns

```
[127]: c_connections = connections.rename(columns={
    "c.katana": "facebook",
    "c.android.chrome": "chrome",
    "c.android.gm": "gmail",
    "c.dogalize": "dogalize",
    "c.android.youtube": "youtube",
    "c.updateassist": "updateassist",
    "c.UCMobile.intl": "UCMobile.intl",
    "c.raider": "raider",
    "c.android.vending": "vending",
    "c.UCMobile.x86": "UCMobile.x86",
})

p_processes = processes.rename(columns={
    "p.katana": "facebook",
    "p.android.chrome": "chrome",
    "p.android.gm": "gmail",
    "p.dogalize": "dogalize",
    "p.android.vending": "vending",
    "p.android.packageinstaller": "packageinstaller",
    "p.system": "system",
    "p.android.documentsui": "documentsui",
    "p.android.settings": "settings",
    "p.android.externalstorage": "externalstorage",
    "p.android.defcontainer": "defcontainer",
    "p.inputmethod.latin": "inputmethod.latin",
    "p.process.gapps": "gapps",
    "p.simulator": "simulator",
    "p.android.gms": "google mobile services (gms)",
    "p.google": "google",
    "p.olauncher": "olauncher",
    "p.browser.provider": "browser provider",
    "p.notifier": "notifier",
    "p.gms.persistent": "gms.persistent",
})
```

```
} )
```

Change the type of timestamp to int64 of connections' dataset:

```
[128]: c_connections['ts'] = pd.to_datetime(c_connections['ts']).astype(np.int64)
```

Change the type of timestamp to int64 of processes dataset:

```
[129]: p_processes['ts'] = pd.to_datetime(p_processes['ts']).astype(np.int64)
```

1.1.1 Merge datasets connections and processes:

```
[130]: merged_dataset = pd.merge(c_connections, p_processes, on=['imei', 'ts', 'mwra'])
merged_dataset
```

```
[130]:
```

		ts	imei	mwra	facebook_x	chrome_x	\
0		1525514400000000000	3590433799317662188	1.0	10.99774	11.05477	
1		1525514460000000000	3590433799317662394	1.0	11.08234	9.64636	
2		1525514520000000000	3590433799317661834	0.0	11.49582	12.27416	
3		1525514580000000000	8630330696303481289	0.0	10.50935	11.41774	
4		1525514640000000000	8630330696303481149	0.0	10.25989	14.46448	
...		
15445		1525987200000000000	3590433799317661925	1.0	11.23638	12.54494	
15446		1526140320000000000	863033069630348776	0.0	11.71795	13.86245	
15447		1526140320000000000	863033069630348776	0.0	11.71795	13.86245	
15448		1526338140000000000	3590433799317661206	1.0	15.51197	16.53309	
15449		1526338140000000000	3590433799317661206	1.0	15.51197	16.53309	
	gmail_x	dogalize_x	youtube	updateassist	UCMobile.intl	...	\
0	6.03999	12.49767	8.59956	14.00953	52.54470	...	
1	8.64167	12.60788	9.84197	38.27736	44.56009	...	
2	11.59681	12.99258	9.74923	57.41411	36.83333	...	
3	14.43350	12.91018	13.93857	31.57549	41.34296	...	
4	14.02728	8.58832	13.04853	49.47100	38.86755	...	
...	
15445	11.26646	9.03636	12.76080	40.59272	30.60551	...	
15446	12.07446	10.10313	13.96660	45.07102	52.21975	...	
15447	12.07446	10.10313	13.96660	45.07102	52.21975	...	
15448	15.75924	11.99555	14.99913	43.95935	35.75699	...	
15449	15.75924	11.99555	14.99913	43.95935	35.75699	...	
	dogalize_y	gapps	simulator	facebook_y	\		
0	95.23250	99.55387	82.64951	55.62534			
1	73.67809	55.93619	27.33158	68.28812			
2	49.43847	92.96630	54.04233	25.01599			
3	71.37356	8.34277	87.09809	5.21806			
4	14.58892	27.72954	81.20459	22.42807			
...			

15445	87.48802	1.93764	55.19853	61.24749
15446	54.55465	34.82606	68.63334	98.60112
15447	54.55465	34.82606	68.63334	98.60112
15448	74.14262	50.12292	55.73990	82.29958
15449	74.14262	50.12292	55.73990	82.29958

	google mobile services (gms)	google	olauncher	browser	provider \
0	43.73958	28.79282	8.22474		73.26391
1	67.18486	19.40350	19.26265		58.69464
2	57.15110	60.38043	16.88231		55.62452
3	98.58641	97.22889	37.30215		68.75315
4	25.06680	73.26831	43.72205		78.80356
...	
15445	27.68095	7.42961	73.85143		98.09352
15446	20.39117	17.48535	97.03999		59.31267
15447	20.39117	17.48535	97.03999		59.31267
15448	34.64708	85.38809	68.57332		50.53580
15449	34.64708	85.38809	68.57332		50.53580

	notifier	gms.persistent
0	25.28004	86.66346
1	90.54099	33.10194
2	16.82005	81.58652
3	26.44336	79.98101
4	16.55350	75.03307
...
15445	39.51722	29.41750
15446	97.53013	7.77545
15447	97.53013	7.77545
15448	61.50573	81.83581
15449	61.50573	81.83581

[15450 rows x 33 columns]

1.1.2 Check for missing values and duplicities

```
[131]: has_nan = merged_dataset.isnull().values.any()

if has_nan:
    print("The dataset has NaN values.")
    print(merged_dataset.isnull().values)
else:
    print("No NaN values found in the dataset.")
```

No NaN values found in the dataset.

```
[132]: has_duplicity = merged_dataset.duplicated().any()

if has_duplicity:
    print("The dataset has duplicity values.")
    print(merged_dataset[merged_dataset.duplicated()])
    print("Number of duplicate rows:", merged_dataset.duplicated().sum())
else:
    print("No duplicity values found in the dataset.")
```

The dataset has duplicity values.

	ts	imei	mwra	facebook_x	chrome_x	\
95	15255200400000000000	8630330696303482303	1.0	10.82916	11.29582	
108	15255207600000000000	863033069630348065	1.0	14.71992	14.01692	
175	15255247200000000000	359043379931766353	0.0	8.46728	13.41079	
300	15255321600000000000	8630330696303481669	1.0	9.87679	8.52849	
303	15255322800000000000	359043379931766924	1.0	10.22849	11.46160	
...	
15445	15259872000000000000	3590433799317661925	1.0	11.23638	12.54494	
15446	15261403200000000000	863033069630348776	0.0	11.71795	13.86245	
15447	15261403200000000000	863033069630348776	0.0	11.71795	13.86245	
15448	15263381400000000000	3590433799317661206	1.0	15.51197	16.53309	
15449	15263381400000000000	3590433799317661206	1.0	15.51197	16.53309	

	gmail_x	dogalize_x	youtube	updateassist	UCMobile.intl	...	\
95	12.66003	10.16379	16.04734	39.11921	41.94816	...	
108	14.46679	10.37614	15.28242	54.88030	54.02226	...	
175	11.08708	6.18605	9.60942	52.78819	56.42743	...	
300	9.57042	11.81096	7.54489	7.94035	61.92247	...	
303	8.69223	10.45511	12.44206	46.64713	52.17634	...	
...	
15445	11.26646	9.03636	12.76080	40.59272	30.60551	...	
15446	12.07446	10.10313	13.96660	45.07102	52.21975	...	
15447	12.07446	10.10313	13.96660	45.07102	52.21975	...	
15448	15.75924	11.99555	14.99913	43.95935	35.75699	...	
15449	15.75924	11.99555	14.99913	43.95935	35.75699	...	

	dogalize_y	gapps	simulator	facebook_y	\
95	97.52469	95.01733	56.82875	81.39681	
108	87.62723	75.66986	13.90231	10.51647	
175	30.91028	1.84816	36.52292	8.24216	
300	75.93998	80.95502	86.96990	1.35263	
303	61.25695	57.35034	36.31488	67.14928	
...	
15445	87.48802	1.93764	55.19853	61.24749	
15446	54.55465	34.82606	68.63334	98.60112	
15447	54.55465	34.82606	68.63334	98.60112	
15448	74.14262	50.12292	55.73990	82.29958	
15449	74.14262	50.12292	55.73990	82.29958	

	google mobile services (gms)	google	olauncher	browser	provider \
95	15.49079	33.46497	47.35665		9.96695
108	36.08753	32.65558	13.72486		74.68378
175	85.18883	85.80315	28.04099		5.37063
300	33.50430	7.26696	23.43879		90.34312
303	38.96647	56.28646	88.96592		11.03558
...
15445	27.68095	7.42961	73.85143		98.09352
15446	20.39117	17.48535	97.03999		59.31267
15447	20.39117	17.48535	97.03999		59.31267
15448	34.64708	85.38809	68.57332		50.53580
15449	34.64708	85.38809	68.57332		50.53580

	notifier	gms.persistent
95	9.45719	90.77472
108	47.71853	10.51180
175	64.96796	16.24571
300	43.77779	19.04622
303	36.52814	89.92218
...
15445	39.51722	29.41750
15446	97.53013	7.77545
15447	97.53013	7.77545
15448	61.50573	81.83581
15449	61.50573	81.83581

[537 rows x 33 columns]

Number of duplicate rows: 537

1.1.3 Drop values which are not helpful for further training:

```
[133]: merged_dataset.drop('ts', axis=1, inplace=True)
merged_dataset.drop('imei', axis=1, inplace=True)
```

1.1.4 Outlier deletion

```
[134]: # Source: https://www.kaggle.com/code/marcinrutecki/outlier-detection-methods

def StandardDevDetection(data, n, columns):

    outliers_inx = []
    lower = 0
    upper = 0

    for column in columns:
        # Calculate mean and standard derivation of each column
```

```

        data_mean, data_std = mean(data[column], axis=0), std(data[column],
↪axis=0)
        print('column=', column, 'len=', len(data), 'mean=', data_mean, 'std=',
↪data_std)

        # Divide it to the three outliers in the standard deviations:
        cut_off = data_std * 3
        lower, upper = data_mean - cut_off, data_mean + cut_off
        print('column=', column, 'cutoff=', cut_off, 'lower=', lower, 'upper=',
↪upper)

        # Filter the dataframe:
        outliers = data[(data[column] < lower) | (data[column] > upper)].index
        print('Identified outliers:', len(outliers))

        outliers_inx.extend(outliers)

        outliers_inx = Counter(outliers_inx)
        multiple_outliers = list( k for k, v in outliers_inx.items() if v > n )

        data_upper = data[data[column] > upper]
        data_lower = data[data[column] < lower]
        print('Total number of outliers is:', data_upper.shape[0] + data_lower.
↪shape[0])

        return multiple_outliers

columns = merged_dataset.columns
result = StandardDevDetection(merged_dataset, 1, columns)

new_dataset = merged_dataset.drop(result, axis = 0).reset_index(drop=True)

```

```

column= mwra len= 15450 mean= 0.6255663430420711 std= 0.48397633567669496
column= mwra cutoff= 1.4519290070300848 lower= -0.8263626639880136 upper=
2.077495350072156
Identified outliers: 0
column= facebook_x len= 15450 mean= 10.962643722330096 std= 2.6723458605226362
column= facebook_x cutoff= 8.017037581567909 lower= 2.9456061407621874 upper=
18.979681303898005
Identified outliers: 30
column= chrome_x len= 15450 mean= 11.605878082200647 std= 2.5788799704504304
column= chrome_x cutoff= 7.736639911351292 lower= 3.8692381708493553 upper=
19.342517993551937
Identified outliers: 20
column= gmail_x len= 15450 mean= 12.255588944983819 std= 2.5607912382789277
column= gmail_x cutoff= 7.682373714836784 lower= 4.573215230147035 upper=

```

19.937962659820602
 Identified outliers: 47
 column= dogalize_x len= 15450 mean= 10.453980822653723 std= 2.2933559796410883
 column= dogalize_x cutoff= 6.880067938923265 lower= 3.5739128837304577 upper= 17.334048761576987
 Identified outliers: 79
 column= youtube len= 15450 mean= 12.250241922330098 std= 2.55734941631557
 column= youtube cutoff= 7.67204824894671 lower= 4.578193673383388 upper= 19.922290171276806
 Identified outliers: 14
 column= updateassist len= 15450 mean= 45.98381515598705 std= 12.502608454363255
 column= updateassist cutoff= 37.507825363089765 lower= 8.475989792897288 upper= 83.49164051907681
 Identified outliers: 42
 column= UCMobile.intl len= 15450 mean= 45.88107211326861 std= 13.050892012462258
 column= UCMobile.intl cutoff= 39.152676037386776 lower= 6.728396075881832 upper= 85.03374815065538
 Identified outliers: 24
 column= raider len= 15450 mean= 49.188615063430426 std= 13.337020785059845
 column= raider cutoff= 40.01106235517953 lower= 9.177552708250893 upper= 89.19967741860995
 Identified outliers: 38
 column= vending_x len= 15450 mean= 49.60707255469255 std= 28.92403819144008
 column= vending_x cutoff= 86.77211457432024 lower= -37.16504201962769 upper= 136.3791871290128
 Identified outliers: 0
 column= UCMobile.x86 len= 15450 mean= 49.76764124854368 std= 28.696862484297487
 column= UCMobile.x86 cutoff= 86.09058745289246 lower= -36.32294620434878 upper= 135.85822870143613
 Identified outliers: 0
 column= packageinstaller len= 15450 mean= 11.078107833009707 std= 2.814283167101723
 column= packageinstaller cutoff= 8.442849501305169 lower= 2.635258331704538 upper= 19.520957334314875
 Identified outliers: 22
 column= system len= 15450 mean= 10.999383213592232 std= 2.5257279118949145
 column= system cutoff= 7.577183735684743 lower= 3.4221994779074887 upper= 18.576566949276973
 Identified outliers: 26
 column= documentsui len= 15450 mean= 11.08572180064725 std= 2.5594279088864607
 column= documentsui cutoff= 7.678283726659382 lower= 3.4074380739878674 upper= 18.76400552730663
 Identified outliers: 23
 column= chrome_y len= 15450 mean= 12.138836054368932 std= 2.5379252214438917
 column= chrome_y cutoff= 7.6137756643316745 lower= 4.5250603900372575 upper= 19.752611718700607
 Identified outliers: 22
 column= settings len= 15450 mean= 13.40750721618123 std= 1.898547862388675

column= settings cutoff= 5.695643587166025 lower= 7.711863629015205 upper= 19.103150803347255
 Identified outliers: 48
 column= gmail_y len= 15450 mean= 12.784054436245954 std= 2.4946624687029955
 column= gmail_y cutoff= 7.483987406108986 lower= 5.300067030136968 upper= 20.26804184235494
 Identified outliers: 104
 column= externalstorage len= 15450 mean= 11.600245935922329 std= 2.586746058468083
 column= externalstorage cutoff= 7.760238175404249 lower= 3.8400077605180805 upper= 19.360484111326578
 Identified outliers: 19
 column= defcontainer len= 15450 mean= 50.635713223300975 std= 12.747023639153177
 column= defcontainer cutoff= 38.24107091745953 lower= 12.394642305841444 upper= 88.8767841407605
 Identified outliers: 38
 column= vending_y len= 15450 mean= 0.12987212038834953 std= 1.3552755036182966
 column= vending_y cutoff= 4.06582651085489 lower= -3.9359543904665406 upper= 4.19569863124324
 Identified outliers: 61
 column= inputmethod.latin len= 15450 mean= 50.83265873851133 std= 13.056312844657608
 column= inputmethod.latin cutoff= 39.16893853397282 lower= 11.663720204538507 upper= 90.00159727248415
 Identified outliers: 34
 column= dogalize_y len= 15450 mean= 49.48587180970874 std= 28.911527705274644
 column= dogalize_y cutoff= 86.73458311582394 lower= -37.2487113061152 upper= 136.2204549255327
 Identified outliers: 0
 column= gapps len= 15450 mean= 49.9632538355987 std= 28.846959030294013
 column= gapps cutoff= 86.54087709088203 lower= -36.577623255283335 upper= 136.50413092648074
 Identified outliers: 0
 column= simulator len= 15450 mean= 49.75620074951457 std= 28.925657829696828
 column= simulator cutoff= 86.77697348909048 lower= -37.02077273957591 upper= 136.53317423860506
 Identified outliers: 0
 column= facebook_y len= 15450 mean= 49.80526438187702 std= 28.98326890126
 column= facebook_y cutoff= 86.94980670378 lower= -37.144542321902975 upper= 136.75507108565702
 Identified outliers: 0
 column= google mobile services (gms) len= 15450 mean= 50.24647446407767 std= 28.84846131763429
 column= google mobile services (gms) cutoff= 86.54538395290287 lower= -36.298909488825196 upper= 136.79185841698055
 Identified outliers: 0
 column= google len= 15450 mean= 50.34342666990291 std= 28.79696748375251
 column= google cutoff= 86.39090245125753 lower= -36.04747578135462 upper=

```

136.73432912116044
Identified outliers: 0
column= olauncher len= 15450 mean= 49.943178990291266 std= 29.069115788777484
column= olauncher cutoff= 87.20734736633246 lower= -37.26416837604119 upper=
137.15052635662371
Identified outliers: 0
column= browser provider len= 15450 mean= 49.762283370226534 std=
28.883676043571135
column= browser provider cutoff= 86.6510281307134 lower= -36.888744760486865
upper= 136.41331150093993
Identified outliers: 0
column= notifier len= 15450 mean= 49.63440627572815 std= 29.043983532101446
column= notifier cutoff= 87.13195059630434 lower= -37.49754432057618 upper=
136.76635687203247
Identified outliers: 0
column= gms.persistent len= 15450 mean= 49.77950399288026 std=
28.813966764780407
column= gms.persistent cutoff= 86.44190029434122 lower= -36.662396301460966
upper= 136.22140428722147
Identified outliers: 0
Total number of outliers is: 0

```

1.2 Data splitting

```

[135]: mwra = new_dataset[target_column]

data = new_dataset.drop(columns=[target_column], axis=1)

train_data, test_data, train_mwra, test_mwra = train_test_split(data, mwra,
↳test_size=test_size, random_state=random_state)

[136]: # features selected from our previous analysis
selected_features = ['gmail_x', 'gapps', 'facebook_x', 'chrome_x', 'vending_x',
                    'youtube', 'dogalize_x', 'updateassist', 'UCMobile.intl']

# this class was created with help from Claude.ai
# we were unsure of how to work with pipeline

class DataPreprocessor:
    def __init__(self):
        self.pipeline = None

    def create_pipeline(self):
        numeric_pipeline = Pipeline([
            ('standard_scaler', StandardScaler()),
            ('power_transform', PowerTransformer(method='yeo-johnson')),
            ('minmax_scaler', MinMaxScaler()),

```

```

        ('feature_select', SelectKBest(score_func=f_classif,
↪k=len(selected_features)))
    ])

    self.pipeline = ColumnTransformer(
        transformers=[
            ('numeric', numeric_pipeline, selected_features)
        ],
        remainder='drop' # drop any columns not specified in the
↪transformers
    )

    return self

    def fit_transform(self, X, y=None):
        if self.pipeline is None:
            self.create_pipeline()
        return self.pipeline.fit_transform(X, y)

    def transform(self, X):
        if self.pipeline is None:
            raise ValueError("Pipeline has not been fitted yet. Call
↪fit_transform first.")
        return self.pipeline.transform(X)

def process_data(train_data, test_data, train_mwra=None):
    preprocessor = DataPreprocessor()

    X_train_processed = preprocessor.fit_transform(train_data, train_mwra)
    X_test_processed = preprocessor.transform(test_data)

    return X_train_processed, X_test_processed

```

```

[137]: X_train_processed, X_test_processed = process_data(train_data, test_data,
↪train_mwra)

X_train_processed_df = pd.DataFrame(X_train_processed,
↪columns=selected_features)
X_test_processed_df = pd.DataFrame(X_test_processed, columns=selected_features)

X_train_processed_df.to_csv('dataset_df_train.csv', index=False)
X_test_processed_df.to_csv('dataset_df_test.csv', index=False)

print("Processed training data shape:", X_train_processed_df.shape)
print("Processed test data shape:", X_test_processed_df.shape)
X_train_processed_df

```

Processed training data shape: (10777, 9)

Processed test data shape: (4620, 9)

```
[137]:      gmail_x      gapps  facebook_x  chrome_x  vending_x  youtube  \
0      0.427075  0.444816   0.341966  0.369723   0.510712  0.713017
1      0.781099  0.141603   0.592307  0.357581   0.626303  0.529742
2      0.482209  0.519524   0.410873  0.534186   0.501690  0.476516
3      0.547636  0.015084   0.500967  0.251712   0.541367  0.497678
4      0.319781  0.142447   0.574914  0.393006   0.267689  0.286099
...      ...      ...      ...      ...      ...      ...
10772  0.555393  0.700282   0.636286  0.557965   0.745899  0.342406
10773  0.483965  0.486749   0.683975  0.600363   0.779375  0.473029
10774  0.404798  0.737832   0.420983  0.499295   0.858299  0.407645
10775  0.623811  0.926622   0.708997  0.666297   0.713649  0.478073
10776  0.258322  0.887967   0.614210  0.506339   0.229566  0.337356

      dogalize_x  updateassist  UCMobile.intl
0      0.462090      0.421345      0.266698
1      0.213026      0.617507      0.531483
2      0.460207      0.457168      0.498022
3      0.632250      0.246029      0.593191
4      0.442534      0.695363      0.430522
...      ...      ...      ...
10772  0.672724      0.155965      0.680182
10773  0.567626      0.391143      0.318646
10774  0.340618      0.520895      0.523320
10775  0.510748      0.420184      0.280551
10776  0.380861      0.324609      0.396991
```

[10777 rows x 9 columns]